

Fault Tolerance



Fault Tolerance

- ▶ **Fault tolerance** is the ability of a distributed system to provide its services even in the presence of faults.

Fault Tolerance

- ▶ **Fault tolerance** is the ability of a distributed system to provide its services even in the presence of faults.
- ▶ A distributed system should be able to recover automatically from partial failures without seriously affecting availability and performance.

Fault Tolerance: Basic Concepts (1)

- ▶ Being fault tolerant is strongly related to what are called dependable systems **Dependability** implies the following:

Fault Tolerance: Basic Concepts (1)

- ▶ Being fault tolerant is strongly related to what are called dependable systems **Dependability** implies the following:
 - ▶ Availability
 - ▶ Reliability
 - ▶ Safety
 - ▶ Maintainability

Fault Tolerance: Basic Concepts (1)

- ▶ Being fault tolerant is strongly related to what are called dependable systems **Dependability** implies the following:
 - ▶ Availability
 - ▶ Reliability
 - ▶ Safety
 - ▶ Maintainability
- ▶ **In-class Exercise:** How is availability different from reliability? How about a system that goes down for 1 millisecond every hour? How about a system that never goes down but has to be shut down two weeks every year?

Fault Tolerance: Basic Concepts (2)

- ▶ A system is said to **fail** when it cannot meet its promises. For example, if it cannot provide a service (or only provide it partially).

Fault Tolerance: Basic Concepts (2)

- ▶ A system is said to **fail** when it cannot meet its promises. For example, if it cannot provide a service (or only provide it partially).
- ▶ An **error** is a part of the system's state that may lead to failure. The cause of an error is called a **fault**.

Fault Tolerance: Basic Concepts (2)

- ▶ A system is said to **fail** when it cannot meet its promises. For example, if it cannot provide a service (or only provide it partially).
- ▶ An **error** is a part of the system's state that may lead to failure. The cause of an error is called a **fault**.
- ▶ Types of faults:
 - ▶ **Transient**: Occurs once and then disappears. If a operation is repeated, the fault goes away.

Fault Tolerance: Basic Concepts (2)

- ▶ A system is said to **fail** when it cannot meet its promises. For example, if it cannot provide a service (or only provide it partially).
- ▶ An **error** is a part of the system's state that may lead to failure. The cause of an error is called a **fault**.
- ▶ Types of faults:
 - ▶ **Transient**: Occurs once and then disappears. If a operation is repeated, the fault goes away.
 - ▶ **Intermittent**: Occurs, then vanishes of its own accord, then reappears and so on.

Fault Tolerance: Basic Concepts (2)

- ▶ A system is said to **fail** when it cannot meet its promises. For example, if it cannot provide a service (or only provide it partially).
- ▶ An **error** is a part of the system's state that may lead to failure. The cause of an error is called a **fault**.
- ▶ Types of faults:
 - ▶ **Transient**: Occurs once and then disappears. If a operation is repeated, the fault goes away.
 - ▶ **Intermittent**: Occurs, then vanishes of its own accord, then reappears and so on.
 - ▶ **Permanent**: Continues to exist until the faulty component is replaced.

Fault Tolerance: Basic Concepts (2)

- ▶ A system is said to **fail** when it cannot meet its promises. For example, if it cannot provide a service (or only provide it partially).
- ▶ An **error** is a part of the system's state that may lead to failure. The cause of an error is called a **fault**.
- ▶ Types of faults:
 - ▶ **Transient**: Occurs once and then disappears. If a operation is repeated, the fault goes away.
 - ▶ **Intermittent**: Occurs, then vanishes of its own accord, then reappears and so on.
 - ▶ **Permanent**: Continues to exist until the faulty component is replaced.
- ▶ **In-class Exercise**: Give examples of each type of fault for your project!

Failure Models

| Type of failure | Description |
|---|--|
| Crash failure | A server halts, but is working correctly until it halts |
| Omission failure <i>Receive omission</i> <i>Send omission</i> | A server fails to respond to incoming requests A server fails to receive incoming messages A server fails to send messages |
| Timing failure | A server's response lies outside the specified time interval |
| Response failure <i>Value failure</i> <i>State transition failure</i> | A server's response is incorrect The value of the response is wrong The server deviates from the correct flow of control |
| Arbitrary failure | A server may produce arbitrary responses at arbitrary times |

Failure Models

| Type of failure | Description |
|---|--|
| Crash failure | A server halts, but is working correctly until it halts |
| Omission failure <i>Receive omission</i> <i>Send omission</i> | A server fails to respond to incoming requests A server fails to receive incoming messages A server fails to send messages |
| Timing failure | A server's response lies outside the specified time interval |
| Response failure <i>Value failure</i> <i>State transition failure</i> | A server's response is incorrect The value of the response is wrong The server deviates from the correct flow of control |
| Arbitrary failure | A server may produce arbitrary responses at arbitrary times |

- **Fail-stop** versus **fail-silent**.

Failure Models

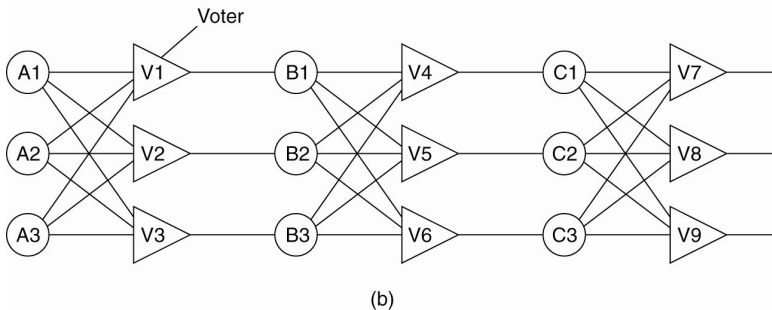

| Type of failure | Description |
|---|--|
| Crash failure | A server halts, but is working correctly until it halts |
| Omission failure <i>Receive omission</i> <i>Send omission</i> | A server fails to respond to incoming requests A server fails to receive incoming messages A server fails to send messages |
| Timing failure | A server's response lies outside the specified time interval |
| Response failure <i>Value failure</i> <i>State transition failure</i> | A server's response is incorrect The value of the response is wrong The server deviates from the correct flow of control |
| Arbitrary failure | A server may produce arbitrary responses at arbitrary times |

- ▶ **Fail-stop** versus **fail-silent**.
- ▶ **Byzantine Failures:**
 - ▶ Arbitrary failures where a server is producing output that it should never have produced, but which cannot be detected as being incorrect.
 - ▶ A faulty server may even be working with other servers to produce intentionally wrong answers!

Failure Masking by Redundancy

- ▶ **Information Redundancy.** For example, adding extra bits (like in Hamming Codes, see the book Coding and Information Theory) to allow recovery from garbled bits.
- ▶ **Time Redundancy.** Repeat actions if need be.
- ▶ **Physical Redundancy.** Extra equipment or processes are added to make the system tolerate loss of some components.



Failure Masking by Physical Redundancy



Achieved by replicating processes into groups.

- ▶ How to design fault-tolerant groups?
- ▶ How to reach an agreement within a group when some members cannot be trusted to give correct answers?

Flat Groups Versus Hierarchical Groups



Failure Masking Via Replication

- ▶ *Primary-backup protocol*. A primary coordinates all write operations. If it fails, then the others hold an election to replace the primary
- ▶ *Replicated-write protocols*. Active replication as well as quorum based protocols. Corresponds to a flat group

Failure Masking Via Replication

- ▶ *Primary-backup protocol*. A primary coordinates all write operations. If it fails, then the others hold an election to replace the primary
- ▶ *Replicated-write protocols*. Active replication as well as quorum based protocols. Corresponds to a flat group
- ▶ A system is said to be **k fault tolerant** if it can survive faults in k components and still meet its specifications.

Failure Masking Via Replication

- ▶ *Primary-backup protocol*. A primary coordinates all write operations. If it fails, then the others hold an election to replace the primary
- ▶ *Replicated-write protocols*. Active replication as well as quorum based protocols. Corresponds to a flat group
- ▶ A system is said to be **k fault tolerant** if it can survive faults in k components and still meet its specifications.
 - ▶ For *fail-silent* components, $k + 1$ are enough to be k fault tolerant.

Failure Masking Via Replication

- ▶ *Primary-backup protocol*. A primary coordinates all write operations. If it fails, then the others hold an election to replace the primary
- ▶ *Replicated-write protocols*. Active replication as well as quorum based protocols. Corresponds to a flat group
- ▶ A system is said to be **k fault tolerant** if it can survive faults in k components and still meet its specifications.
 - ▶ For *fail-silent* components, $k + 1$ are enough to be k fault tolerant.
 - ▶ For *Byzantine failures*, at least $2k + 1$ extra components are needed to achieve k fault tolerance.

Failure Masking Via Replication

- ▶ *Primary-backup protocol*. A primary coordinates all write operations. If it fails, then the others hold an election to replace the primary
- ▶ *Replicated-write protocols*. Active replication as well as quorum based protocols. Corresponds to a flat group
- ▶ A system is said to be **k fault tolerant** if it can survive faults in k components and still meet its specifications.
 - ▶ For *fail-silent* components, $k + 1$ are enough to be k fault tolerant.
 - ▶ For *Byzantine failures*, at least $2k + 1$ extra components are needed to achieve k fault tolerance.
 - ▶ Requires atomic multicasting: all requests arrive at all servers in same order. This can be relaxed to just for write operations.

Agreement in Faulty Systems (1)

The general goal of **distributed agreement algorithms** is to have all the non-faulty processes reach agree consensus on some issue, and to establish that consensus within a finite number of steps.

Agreement in Faulty Systems (1)


The general goal of **distributed agreement algorithms** is to have all the non-faulty processes reach agree consensus on some issue, and to establish that consensus within a finite number of steps. Possible cases:

- ▶ Synchronous versus asynchronous systems
- ▶ Communication delay is bounded or not
- ▶ Message delivery is ordered or not
- ▶ Message transmission is done through unicasting or multicasting

Agreement in Faulty Systems (2)

| | | Message ordering | | | | Communication delay |
|------------------|--------------|------------------|-----------|---------|-----------|---------------------|
| | | Unordered | | Ordered | | |
| | | Unicast | Multicast | Unicast | Multicast | |
| Process behavior | Synchronous | | | X | | Bounded |
| | | | | X | | Unbounded |
| | Asynchronous | X | X | X | X | Bounded |
| | | | | X | X | Unbounded |

Agreement in Faulty Systems (3)



- ▶ Two Army Problem
- ▶ Non-faulty generals with unreliable communication.

Agreement in Faulty Systems (3)

Byzantine Generals Problem (Lamport)

- ▶ Red army in the valley, n blue generals each with their own army surrounding them.
- ▶ Communication is pairwise, instantaneous and perfect.
- ▶ However m of the blue generals are traitors (faulty processes) and are actively trying to prevent the loyal generals from reaching agreement. The generals know the value m .

Agreement in Faulty Systems (3)

Byzantine Generals Problem (Lamport)

- ▶ Red army in the valley, n blue generals each with their own army surrounding them.
- ▶ Communication is pairwise, instantaneous and perfect.
- ▶ However m of the blue generals are traitors (faulty processes) and are actively trying to prevent the loyal generals from reaching agreement. The generals know the value m .

Goal: The generals need to exchange their troop strengths. At the end of the algorithm, each general has a vector of length n . If i th general is loyal, then the i th element has their correct troop strength otherwise it is undefined.

Agreement in Faulty Systems (3)

Byzantine Generals Problem (Lamport)


- ▶ Red army in the valley, n blue generals each with their own army surrounding them.
- ▶ Communication is pairwise, instantaneous and perfect.
- ▶ However m of the blue generals are traitors (faulty processes) and are actively trying to prevent the loyal generals from reaching agreement. The generals know the value m .

Goal: The generals need to exchange their troop strengths. At the end of the algorithm, each general has a vector of length n . If i th general is loyal, then the i th element has their correct troop strength otherwise it is undefined.

Conditions for a solution:

- ▶ All loyal generals decide upon the same plan of action
- ▶ A small number of traitors cannot cause the loyal generals to adopt a bad plan

Byzantine Agreement Example (1)



(a)

1 Got(1, 2, x, 4)
 2 Got(1, 2, y, 4)
 3 Got(1, 2, 3, 4)
 4 Got(1, 2, z, 4)

(b)


| | | |
|--------------|--------------|--------------|
| 1 Got | 2 Got | 4 Got |
| (1, 2, y, 4) | (1, 2, x, 4) | (1, 2, x, 4) |
| (a, b, c, d) | (e, f, g, h) | (1, 2, y, 4) |
| (1, 2, z, 4) | (1, 2, z, 4) | (i, j, k, l) |

(c)

The Byzantine generals problem for 3 loyal generals and 1 traitor:

- ▶ The generals announce their troop strengths (let's say, in units of 1 kilo soldiers)
- ▶ The vectors that each general assembles based on previous step
- ▶ The vectors that each general receives
- ▶ If a value has a majority, then we know it correctly, else it is unknown

Byzantine Agreement Example (2)



(a)

1 Got(1, 2, x)
2 Got(1, 2, y)
3 Got(1, 2, 3)


(b)

| | |
|-----------|-----------|
| 1 Got | 2 Got |
| (1, 2, y) | (1, 2, x) |
| (a, b, c) | (d, e, f) |

(c)

- ▶ The same as in previous slide, except now with 2 loyal generals and one traitor.

Byzantine Agreement Example (2)



(a)

1 Got(1, 2, x)
2 Got(1, 2, y)
3 Got(1, 2, 3)

(b)

| | |
|-----------|-----------|
| 1 Got | 2 Got |
| (1, 2, y) | (1, 2, x) |
| (a, b, c) | (d, e, f) |

(c)

- ▶ The same as in previous slide, except now with 2 loyal generals and one traitor.
- ▶ For m faulty processes, we need a total of $3m + 1$ processes to reach agreement.

RMI semantics in the presence of failures.

RMI semantics in the presence of failures.

- ▶ The client is unable to locate the server.

RMI semantics in the presence of failures.

- ▶ The client is unable to locate the server.
- ▶ The request message from the client to the server is lost.

RMI semantics in the presence of failures.

- ▶ The client is unable to locate the server.
- ▶ The request message from the client to the server is lost.
- ▶ The server crashes after receiving a request.

RMI semantics in the presence of failures.

- ▶ The client is unable to locate the server.
- ▶ The request message from the client to the server is lost.
- ▶ The server crashes after receiving a request.
- ▶ The reply message from the server to the client is lost.

RMI semantics in the presence of failures.

- ▶ The client is unable to locate the server.
- ▶ The request message from the client to the server is lost.
- ▶ The server crashes after receiving a request.
- ▶ The reply message from the server to the client is lost.
- ▶ The client crashes after sending a request.

RPC Semantics in the Presence of Failures

- ▶ RPC failures
 - ▶ **Client cannot locate server:** Raise an exception or send a signal to client leading to loss in transparency

RPC Semantics in the Presence of Failures

- ▶ RPC failures
 - ▶ **Client cannot locate server**: Raise an exception or send a signal to client leading to loss in transparency
 - ▶ **Lost request messages**: Start a timer when sending a request. If timer expires before a reply is received, send the request again. Server would need to detect duplicate requests

RPC Semantics in the Presence of Failures

- ▶ RPC failures
 - ▶ **Client cannot locate server**: Raise an exception or send a signal to client leading to loss in transparency
 - ▶ **Lost request messages**: Start a timer when sending a request. If timer expires before a reply is received, send the request again. Server would need to detect duplicate requests
 - ▶ **Server crashes**: Server crashes before or after executing the request is indistinguishable from the client side...

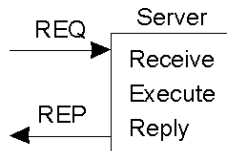
RPC Semantics in the Presence of Failures

- ▶ RPC failures
 - ▶ **Client cannot locate server**: Raise an exception or send a signal to client leading to loss in transparency
 - ▶ **Lost request messages**: Start a timer when sending a request. If timer expires before a reply is received, send the request again. Server would need to detect duplicate requests
 - ▶ **Server crashes**: Server crashes before or after executing the request is indistinguishable from the client side...
- ▶ Possible RPC semantics


RPC Semantics in the Presence of Failures

- ▶ RPC failures
 - ▶ **Client cannot locate server**: Raise an exception or send a signal to client leading to loss in transparency
 - ▶ **Lost request messages**: Start a timer when sending a request. If timer expires before a reply is received, send the request again. Server would need to detect duplicate requests
 - ▶ **Server crashes**: Server crashes before or after executing the request is indistinguishable from the client side...
- ▶ Possible RPC semantics
 - ▶ Exactly once semantics
 - ▶ At least once semantics
 - ▶ At most once semantics
 - ▶ Guarantee nothing semantics


Server Crash (1)



(a)



(b)



(c)

► A server in client-server communication

- (a) Normal case
- (b) Crash after execution
- (c) Crash before execution

Server Crash (2)

- ▶ **Server**: Prints text on receiving request from client and sends message to client after text is printed.

Server Crash (2)

- ▶ **Server:** Prints text on receiving request from client and sends message to client after text is printed.
 - ▶ Send a completion message just before it actually tells the printer to do its work
 - ▶ Or after the text has been printed

Server Crash (2)

- ▶ **Server:** Prints text on receiving request from client and sends message to client after text is printed.
 - ▶ Send a completion message just before it actually tells the printer to do its work
 - ▶ Or after the text has been printed
- ▶ **Client:**

Server Crash (2)

- ▶ **Server:** Prints text on receiving request from client and sends message to client after text is printed.
 - ▶ Send a completion message just before it actually tells the printer to do its work
 - ▶ Or after the text has been printed
- ▶ **Client:**
 - ▶ Never to reissue a request.

Server Crash (2)

- ▶ **Server:** Prints text on receiving request from client and sends message to client after text is printed.
 - ▶ Send a completion message just before it actually tells the printer to do its work
 - ▶ Or after the text has been printed
- ▶ **Client:**
 - ▶ Never to reissue a request.
 - ▶ Always reissue a request.

Server Crash (2)

- ▶ **Server:** Prints text on receiving request from client and sends message to client after text is printed.
 - ▶ Send a completion message just before it actually tells the printer to do its work
 - ▶ Or after the text has been printed
- ▶ **Client:**
 - ▶ Never to reissue a request.
 - ▶ Always reissue a request.
 - ▶ Reissue a request only if it did not receive an acknowledgment of its request being delivered to the server.

Server Crash (2)

- ▶ **Server:** Prints text on receiving request from client and sends message to client after text is printed.
 - ▶ Send a completion message just before it actually tells the printer to do its work
 - ▶ Or after the text has been printed
- ▶ **Client:**
 - ▶ Never to reissue a request.
 - ▶ Always reissue a request.
 - ▶ Reissue a request only if it did not receive an acknowledgment of its request being delivered to the server.
 - ▶ Reissue a request only if it has not received an acknowledgment of its print request.

Server Crash (3)

- ▶ Three events that can happen at the server:
 - ▶ Send the completion message (M)
 - ▶ Print the text (P)
 - ▶ Crash (C)

Server Crash (3)

- ▶ Three events that can happen at the server:
 - ▶ Send the completion message (M)
 - ▶ Print the text (P)
 - ▶ Crash (C)
- ▶ These events can occur in six different orderings:

Server Crash (4)

1. $M \rightarrow P \rightarrow C$: A crash occurs after sending the completion message and printing the text.

Server Crash (4)

1. $M \rightarrow P \rightarrow C$: A crash occurs after sending the completion message and printing the text.
2. $M \rightarrow C(\rightarrow P)$: A crash happens after sending the completion message, but before the text could be printed.

Server Crash (4)

1. $M \rightarrow P \rightarrow C$: A crash occurs after sending the completion message and printing the text.
2. $M \rightarrow C(\rightarrow P)$: A crash happens after sending the completion message, but before the text could be printed.
3. $P \rightarrow M \rightarrow C$: A crash occurs after sending the completion message and printing the text.

Server Crash (4)

1. $M \rightarrow P \rightarrow C$: A crash occurs after sending the completion message and printing the text.
2. $M \rightarrow C(\rightarrow P)$: A crash happens after sending the completion message, but before the text could be printed.
3. $P \rightarrow M \rightarrow C$: A crash occurs after sending the completion message and printing the text.
4. $P \rightarrow C(\rightarrow M)$: The text printed, after which a crash occurs before the completion message could be sent.

Server Crash (4)

1. $M \rightarrow P \rightarrow C$: A crash occurs after sending the completion message and printing the text.
2. $M \rightarrow C(\rightarrow P)$: A crash happens after sending the completion message, but before the text could be printed.
3. $P \rightarrow M \rightarrow C$: A crash occurs after sending the completion message and printing the text.
4. $P \rightarrow C(\rightarrow M)$: The text printed, after which a crash occurs before the completion message could be sent.
5. $C(\rightarrow P \rightarrow M)$: A crash happens before the server could do anything.

Server Crash (4)

1. $M \rightarrow P \rightarrow C$: A crash occurs after sending the completion message and printing the text.
2. $M \rightarrow C(\rightarrow P)$: A crash happens after sending the completion message, but before the text could be printed.
3. $P \rightarrow M \rightarrow C$: A crash occurs after sending the completion message and printing the text.
4. $P \rightarrow C(\rightarrow M)$: The text printed, after which a crash occurs before the completion message could be sent.
5. $C(\rightarrow P \rightarrow M)$: A crash happens before the server could do anything.
6. $C(\rightarrow M \rightarrow P)$: A crash happens before the server could do anything.

Server Crash (5)

Client

Server

Strategy M → P

Strategy P → M

Reissue strategy

| |
|---------------------|
| Always |
| Never |
| Only when ACKed |
| Only when not ACKed |

MPC MC(P) C(MP)

| | | |
|-----|------|------|
| DUP | OK | OK |
| OK | ZERO | ZERO |
| DUP | OK | ZERO |
| OK | ZERO | OK |

PMC PC(M)) C(PM)

| | | |
|-----|-----|------|
| DUP | DUP | OK |
| OK | OK | ZERO |
| DUP | OK | ZERO |
| OK | DUP | OK |

- ▶ *M*: send the completion message
- ▶ *P*: print the text
- ▶ *C*: server crash

RPC Semantics in the Presence of Failures (2)

- ▶ **Lost Reply Messages.** Set a timer on client. If it expires without a reply, then send the request again. If requests are **idempotent**, then they can be repeated again without ill-effects

RPC Semantics in the Presence of Failures (2)

- ▶ **Lost Reply Messages.** Set a timer on client. If it expires without a reply, then send the request again. If requests are **idempotent**, then they can be repeated again without ill-effects
- ▶ **Client Crashes.** Creates orphans. An orphan is an active computation on the server for which there is no client waiting.
How to deal with orphans:

RPC Semantics in the Presence of Failures (2)

- ▶ **Lost Reply Messages**. Set a timer on client. If it expires without a reply, then send the request again. If requests are **idempotent**, then they can be repeated again without ill-effects
- ▶ **Client Crashes**. Creates orphans. An orphan is an active computation on the server for which there is no client waiting.
How to deal with orphans:
 - ▶ **Extermination**. Client logs each request in a file before sending it. After a reboot the file is checked and the orphan is explicitly killed off. Expensive, cannot locate grand-orphans etc.

RPC Semantics in the Presence of Failures (2)

- ▶ **Lost Reply Messages**. Set a timer on client. If it expires without a reply, then send the request again. If requests are **idempotent**, then they can be repeated again without ill-effects
- ▶ **Client Crashes**. Creates orphans. An orphan is an active computation on the server for which there is no client waiting.

How to deal with orphans:

- ▶ **Extermination**. Client logs each request in a file before sending it. After a reboot the file is checked and the orphan is explicitly killed off. Expensive, cannot locate grand-orphans etc.
- ▶ **Reincarnation**. Divide time into sequentially numbered epochs. When a client reboots, it broadcasts a message declaring a new epoch. This allows servers to terminate orphan computations.

RPC Semantics in the Presence of Failures (2)

- ▶ **Lost Reply Messages**. Set a timer on client. If it expires without a reply, then send the request again. If requests are **idempotent**, then they can be repeated again without ill-effects
- ▶ **Client Crashes**. Creates orphans. An orphan is an active computation on the server for which there is no client waiting.

How to deal with orphans:

- ▶ **Extermination**. Client logs each request in a file before sending it. After a reboot the file is checked and the orphan is explicitly killed off. Expensive, cannot locate grand-orphans etc.
- ▶ **Reincarnation**. Divide time into sequentially numbered epochs. When a client reboots, it broadcasts a message declaring a new epoch. This allows servers to terminate orphan computations.
- ▶ **Gentle Reincarnation**. A server tries to locate the owner of orphans before killing the computation.

RPC Semantics in the Presence of Failures (2)

- ▶ **Lost Reply Messages**. Set a timer on client. If it expires without a reply, then send the request again. If requests are **idempotent**, then they can be repeated again without ill-effects
- ▶ **Client Crashes**. Creates orphans. An orphan is an active computation on the server for which there is no client waiting.

How to deal with orphans:

- ▶ **Extermination**. Client logs each request in a file before sending it. After a reboot the file is checked and the orphan is explicitly killed off. Expensive, cannot locate grand-orphans etc.
- ▶ **Reincarnation**. Divide time into sequentially numbered epochs. When a client reboots, it broadcasts a message declaring a new epoch. This allows servers to terminate orphan computations.
- ▶ **Gentle Reincarnation**. A server tries to locate the owner of orphans before killing the computation.
- ▶ **Expiration**. Each RPC is given a quantum of time to finish its job. If it cannot finish, then it asks for another quantum. After a crash, a client need only wait for a quantum to make sure all orphans are gone.

Idempotent Operations

- ▶ An **idempotent** operation is one that can be repeated as often as necessary without any harm being done. E.g. reading a block from a file.

Idempotent Operations

- ▶ An **idempotent** operation is one that can be repeated as often as necessary without any harm being done. E.g. reading a block from a file.
- ▶ In general, try to make RPC/RMI methods be idempotent if possible. If not, it can be dealt with in a couple of ways.

Idempotent Operations

- ▶ An **idempotent** operation is one that can be repeated as often as necessary without any harm being done. E.g. reading a block from a file.
- ▶ In general, try to make RPC/RMI methods be idempotent if possible. If not, it can be dealt with in a couple of ways.
 - ▶ Use a sequence number with each request so server can detect duplicates. But now the server needs to keep state for each client....

Idempotent Operations

- ▶ An **idempotent** operation is one that can be repeated as often as necessary without any harm being done. E.g. reading a block from a file.
- ▶ In general, try to make RPC/RMI methods be idempotent if possible. If not, it can be dealt with in a couple of ways.
 - ▶ Use a sequence number with each request so server can detect duplicates. But now the server needs to keep state for each client....
 - ▶ Have a bit in the message to distinguish between original and duplicate transmission