Chapter 1: Introduction

# What is a Distributed System?

- *A distributed system is a collection of independent computers that appears to its users as a single coherent system.*

Characteristics of a Distributed System (1)

► Collection of autonomous computing elements
  ► Computing elements (or nodes) can be hardware and/or software
  ► No global clock
  ► Group membership is difficult: Open groups versus Closed groups
  ► Often organized as an overlay network: structured, unstructured, peer-to-peer

Characteristics of a Distributed System (2)

▶ Single coherent system
  ▶ Requires transparency: process execution and data storage
  ▶ Often implemented using middleware: a separate layer of software that is logically placed on top of the operating systems of the computers that are part of the distributed system.
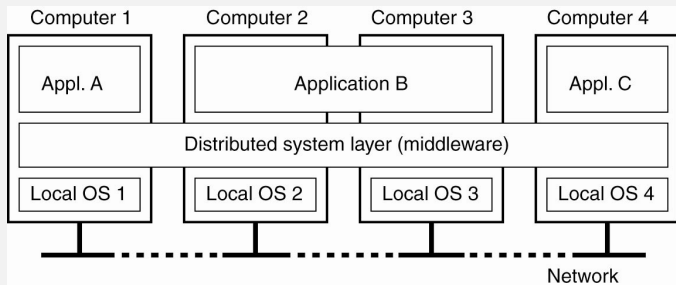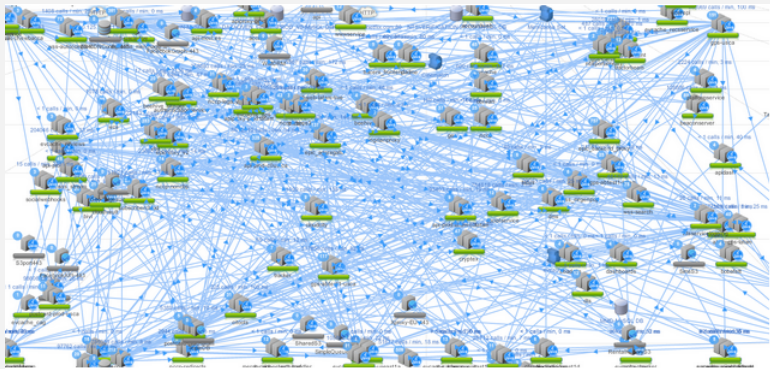

▶ It handles resource management, interapplication communication, masking of and recovery from failures, security and accounting services.

Examples (1)

The Internet is just a world passing around notes in a classroom. –Jon Stewart

Examples (1)

The Internet is just a world passing around notes in a classroom. —Jon Stewart
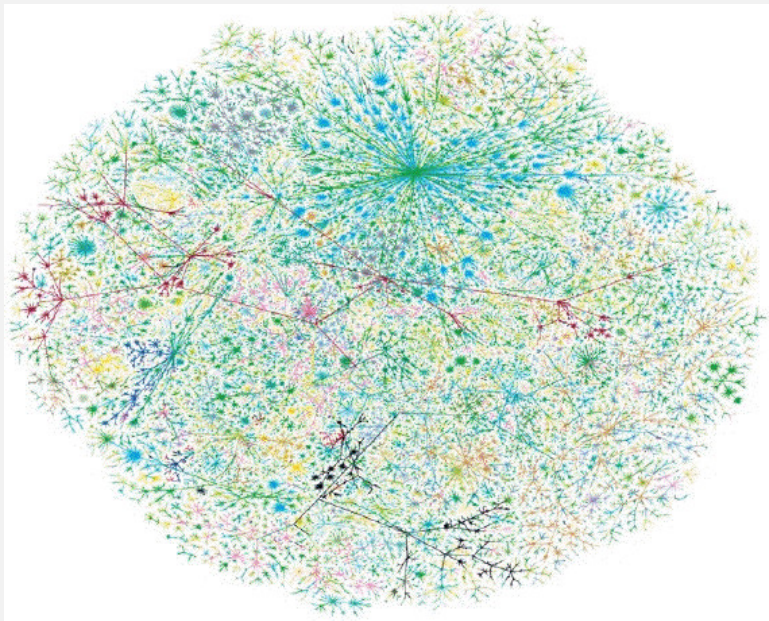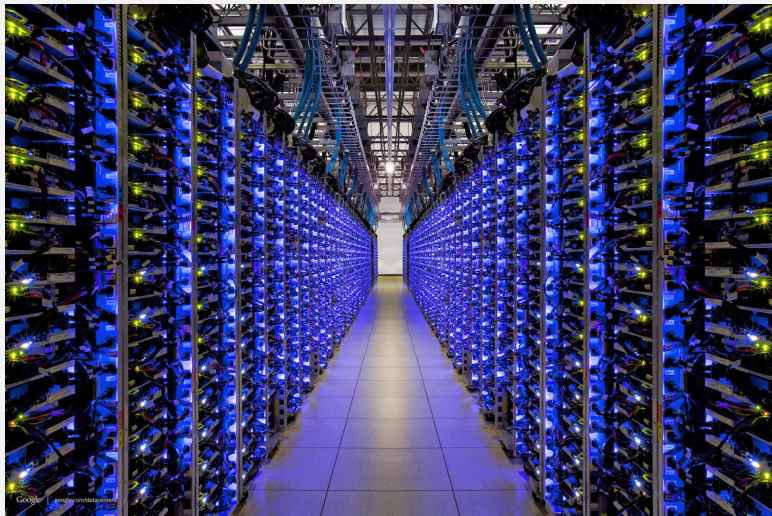
Examples (2)

Examples (2)

Examples (3)

- ▶ Google search, GMail, Google Docs, et al
- ▶ Amazon
- ▶ Cloud storage, such as Dropbox et al
- ▶ Facebook, Twitter, EBay
- ▶ LinkedIn, WhatsApp, Instagram, Snapchat, TikTok
- ▶ Uber, Lyft
- ▶ Hadoop, Hive, Zookeeper, et al
- ▶ No-SQL Databases like HBase, Cassandra, MongoDB, et al
- ▶ Distributed file systems: NFS (Network File System), Google File System (GFS), Hadoop Distributed File System (HDFS), Amazon S3, Lustre, GlusterFS, Parallel Virtual File System (PVFS)
- ▶ P2P: Peer to Peer: Bit Torrent, Gnutella
- ▶ Domain Name System (DNS)
- ▶ Git: distributed version control system.
- ▶ SETI: a distributed computing project in which volunteers donate idle computer power to analyze radio signals for signs of extraterrestrial intelligence.
- ▶ Bitcoin: decentralized digital currency (and other blockchain applications)
- ▶ Virtually every substantial website!

Goals of a Distributed System?

- ▶ Makes resources accessible
- ▶ Reasonably hides the fact that resources are distributed across a network (*Transparency*)
- ▶ Open
- ▶ Scalable

Making Resources Accessible

- Benefits
    - Better economics by sharing expensive resources
    - Easier to collaborate and exchange information
    - Create virtual organizations where geographically dispersed people can work together using groupware
    - Enables electronic commerce
- Problems
    - Eavesdropping or intrusion on communication
    - Tracking of communication to build a profile

Transparency

Type	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may move to another location while in use
Replication	Hide that a resource has multiple copies
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource

Degree of Transparency

Completely hiding the distribution aspects from users is not always a good idea in a distributed system.

- ▶ Attempting to mask a server failure before trying another one may slow down the system
- ▶ Expecting several replicas to be always consistent could degrade performance unacceptably
- ▶ For mobile and embedded devices, it may be better to expose distribution rather than trying to hide it
- ▶ Signal transmission is limited by the speed of light as well as the speed of intermediate switches

Openness

An open distributed system offers services according to standard rules that describe the syntax and semantics of those services.

- ▶ Use of standard protocols.
- ▶ Services are described via interfaces, which are often describe via an Interface Definition Language (IDL). Interfaces only specify syntax so semantics is left to the ambiguities of natural language.
- ▶ Interoperability, Portability, Extensibility.
- ▶ Separating policy from mechanism. For example: *caching* in a web browser.

Scalability

Scalability can be measured against three dimensions.

- ▶ Size: be able to easily add more users and resources to a system
- ▶ Geographical: be able to handle users and resources that are far apart
- ▶ Administrative: be able to manage even if it spans independent administrative organizations

Centralized versus distributed implementations.

Centralized Solutions with Scalability Problems

- ▶ Centralized services.
- ▶ Centralized data.
- ▶ Centralized algorithms.

Distributed Approach

Characteristics of decentralized algorithms:

▶ No machine has complete information about the system state.

▶ Machines make decisions based only on local information.

▶ Failure of one machine does not ruin the algorithm.

▶ There is no implicit assumption that a global clock exists.

In-class exercise. Simulate a centralized and a distributed algorithm for the same problem in class!

Scaling Techniques

- ▶ Hiding communication latencies: Examples would be asynchronous communication as well as pushing code down to clients (E.g. Javascript)
- ▶ Distribution: Taking a component, splitting into smaller parts, and subsequently spreading them across the system. (E.g. Domain Name System)
- ▶ Replication: Replicating components increases availability, helps balance the load leading to better performance, helps hide latencies for geographically distributed systems. Caching is a special form of replication.

Distributed System Development Pitfalls

False assumptions made by first time developer *(formulated by Peter Deustch)*.

- ▶ The network is reliable
- ▶ The network is secure
- ▶ The network is homogeneous
- ▶ The topology does not change
- ▶ Latency is zero
- ▶ Bandwidth is infinite
- ▶ Transport cost is zero
- ▶ There is one administrator

"A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable."
–Leslie Lamport

In-class Exercise

▶ Walk through architecture of various distributed systems ranging from: single server/client, multiple server/clients, point to point and others.

Types of Distributed Systems

- *Distributed Computing Systems*
    - Cluster Computing Systems
    - Grid Computing Systems
    - Cloud Computing
- *Distributed Information Systems*
    - Transaction Processing Systems
    - Enterprise Application Integration
- *Distributed Pervasive Systems*
    - Ubiquitous Computing Systems
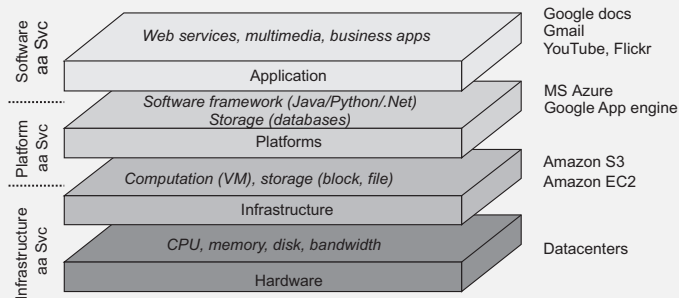   - Mobile Computing Systems
    - Sensor Networks

Cluster Computing

- A computer cluster is primarily used to run a single program in parallel on multiple machines. The program relies on parallel libraries and frameworks such as MPI (Message Passing Interface), MapReduce and others.
- A *computer cluster* consists of a collection of compute and I/O nodes that are controlled and accessed from a single master node. The master allocates nodes to a parallel program, maintains a queue of submitted jobs and provides an interface to manage the cluster.
- A cluster usually has a dedicated high-speed interconnection network and an optional administrative network.
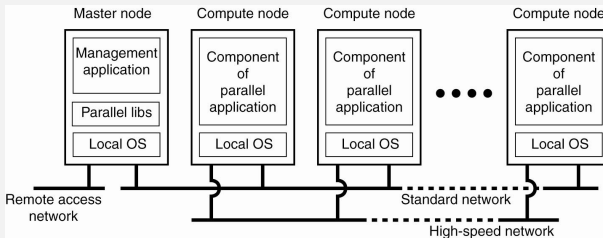
Cloud Computing

▶ Cloud Computing provides the facilities to dynamically construct an infrastructure and compose what is needed from available services.

Layers of Cloud Computing

- ▶ **Hardware**: Processors, routers, power and cooling systems. Customers normally never get to see these.
- ▶ **Infrastructure**: Deploys virtualization techniques. Evolves around allocating and managing virtual storage devices and virtual servers.
- ▶ **Platform**: Provides higher-level abstractions for storage and such. Example: Amazon S3 storage system offers an API for (locally created) files to be organized and stored in so-called buckets.
- ▶ **Application**: Actual applications, whether distributed or not. Includes apps such as office suites (text processors, spreadsheet applications, presentation applications). Comparable to the suite of apps shipped with traditional OSes.

Integrating Applications

- **Situation**: Organizations confronted with many networked applications, but achieving interoperability was painful.

- **Basic approach**: A networked application is one that runs on a server making its services available to remote clients. Simple integration: clients combine requests for (different) applications; send that off; collect responses, and present a coherent result to the user.

- **Next step**: Allow direct application-to-application communication, leading to Enterprise Application Integration (EAI).

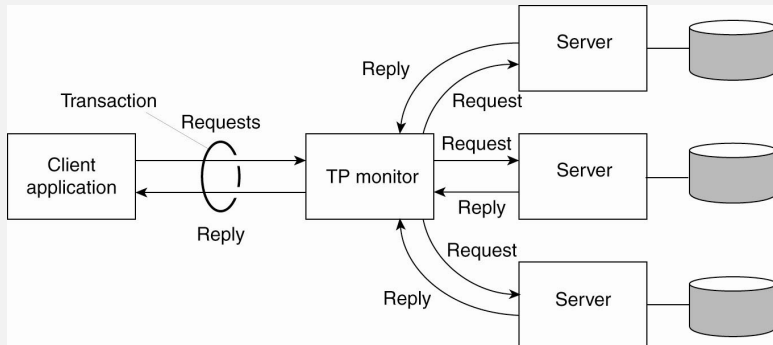# Example EAI: Transaction Processing Systems (1)

Transaction primitives:

- ► BEGIN_TRANSACTION
- ► END_TRANSACTION
- ► ABORT_TRANSACTION
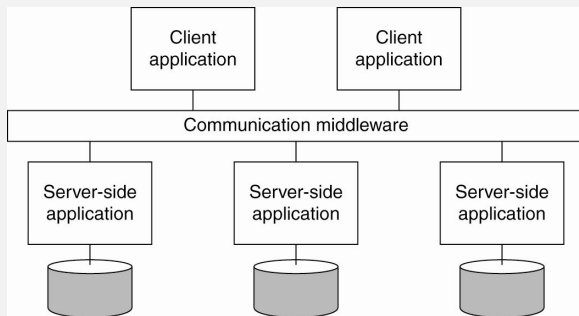- ► READ, WRITE

ACID characteristic properties of transactions:

- ► Atomic: To the outside world, the transaction happens indivisibly
- ► Consistent: The transaction does not violate system invariants
- ► Isolated: Concurrent transactions do not interfere with each other
- ► Durable: Once a transaction commits, the changes are permanent

Transactions can be nested. Durability applies to top-level transactions only. For example: an airline and a hotel database. This led to the design of Transaction Processing Monitors (TP Monitors) to coordinate the commitment of multiple subtransactions.

Enterprise Application Integration

Middleware as a communication facilitator for enterprise application integration. This can be done in multiple ways. For example:

▶ Remote Procedure Call (RPC): Requests are sent through local procedure call, packaged as message, processed, responded through message, and result returned as return from call.

▶ Message Oriented Middleware (MOM): Messages are sent to logical contact point (published), and forwarded to subscribed

Distributed Pervasive Systems

Characteristics of pervasive systems (aka Internet Of Things):

- ▶ Blends into the environment
- ▶ Encourage ad hoc composition
- ▶ Naturally distributed
- ▶ Nodes are often small and battery powered
- ▶ Wireless/mobile communication is the norm

Examples: Ubiquitous Computing Systems, Mobile Computing Systems, Sensor (and actuator) networks

In-class Exercise

In-class Exercise: Classify the distributed system examples we have seen so far into the three categories: Distributed Computing Systems, Distributed Information Systems, Distributed Pervasive Systems.

Chapter 1: Recommended Exercises

- ▶ **Problem 1**. What is the role of middleware in a distributed system?
- ▶ **Problem 2**. Explain what is meant by transparency, and give examples of different types of transparency.
- ▶ **Problem 3**. Scalability can be achieved by applying different techniques. What are these techniques?
- ▶ **Problem 4**. Give further examples of distributed pervasive systems.