


# Process Migration

David Zuercher

CS555 Distributed Systems



# What is Migration?

## ➤ Mi-gra-tion

- The act or an instance of migrating.
- A group migrating together.

## ➤ Mi-gra-te

- Moving from one place to another.

10010001010101  
01101001010101000  
001000100101110  
1101101011100101010  
1010011101001010100010  
001001010101011100101001  
00001010010100010010110101010  
01001001110001010010001011100010

# Why Migrate?

- Some animals migrate for survival.
- A process might do the same if its environment becomes “hostile”:
  - Power outage possible or imminent.
  - The OS is angry?!
  - System becoming overburdened.
- Migration is also fun, and impressive.

# Pseudo-migration

- One can't really *migrate* a process.
- Migration implies a persistence of being (e.g. the PID won't change).
- It is important to acknowledge migration cannot be done literally; but the realization itself is a good starting point for a solution.

point for a solution

# A New Direction

- So, as is often the best solution, we're going to fake it.
- Somehow quantify the state of a process so that a "clone" of it may live on at a later time.


```
101000101
1000101100010
0001001001110001100
1000101110010110110101
011010111010100111010010
10010011100010010101010111
1100000101001010001001011010
1001001001110001010010001011100010
```

```
1001000101010101
0110100101010101000
001000100101110
110110101110010101010
101001110100101010100010
001001010101011100101001
0000101001010001001011010101010
01001001110001010010001011100010
```

# Process Cryonics

- “*Freeze*” a process in its tracks. Save its state into a file of some kind.
- The process state is now fully serialized and static; It could be revived locally or remotely.
- “*Thaw*” the process and let it run as if nothing happened (as close as possible).

# Getting Started: Process Anatomy



- Processes are “physically” composed of memory segments within a virtual address space.
- Segments are mapped one or more pages at a time.
- The heap is usually composed of many segments. The stack is just one that can grow.

0101010101  
0101001010101000  
001000100101110  
110110101110010101010  
101001110100101010100010  
001001010101011100101001  
00001010010100010010110101010  
01001001110001010010001011100010



# The Exact Memory Mapping

In Linux, each process has its own directory under /proc.

A file called "maps" in that directory has the memory mappings for the process.

```
$> cat /proc/<PID>/maps
```

=Address Range	Perm	Offset	Inode	File Mapped=
08048000-08049000	r-xp	00000000	08:04 442116	/home/theprog
08049000-0804b000	rw-p	00000000	08:04 442116	/home/theprog
40000000-40012000	r-xp	00000000	08:02 128783	/lib/ld-2.2.5.so
40012000-40013000	rw-p	00012000	08:02 128783	/lib/ld-2.2.5.so
40013000-40016000	rw-p	00000000	00:00 0	
4002d000-4014b000	r-xp	00000000	08:02 130798	/lib/libc-2.2.5.so
4014b000-40150000	rw-p	0011e000	08:02 130798	/lib/libc-2.2.5.so
40150000-40155000	rw-p	00000000	00:00 0	
bffff3000-c0000000	rwxp	ffff4000	00:00 0	



# What else needs freezing?

- Processes have a “user” area that contains the registers and some other items not accessible to the process itself.
- All variables a process uses are located in the segments mentioned before (so we don't bother with them).
- Per-process Kernel Stack... Let's not worry about that for now.

```
0010001010101
01101001010101000
001000100101110
1101101011100101010
1010011101001010100010
001001010101011100101001
00001010010100010010110101010
01001001110001010010001011100010
```

# Its Getting Colder...

- In Linux, an utterly unrelated process may become the parent of another process (not init!).
- The new parent may “peek” and “poke” one word at a time from the new child’s memory areas.
- Register Data can be peeked and poked as well.

# The Freeze Plan

- I created a *library* function:

```
char *freeze_process(int PID, int ops)
```

- freeze\_process() may be called by one process to stop and freeze another.
- The function returns the name of its primary output file, but there are two...

# \*.freeze & \*.frore

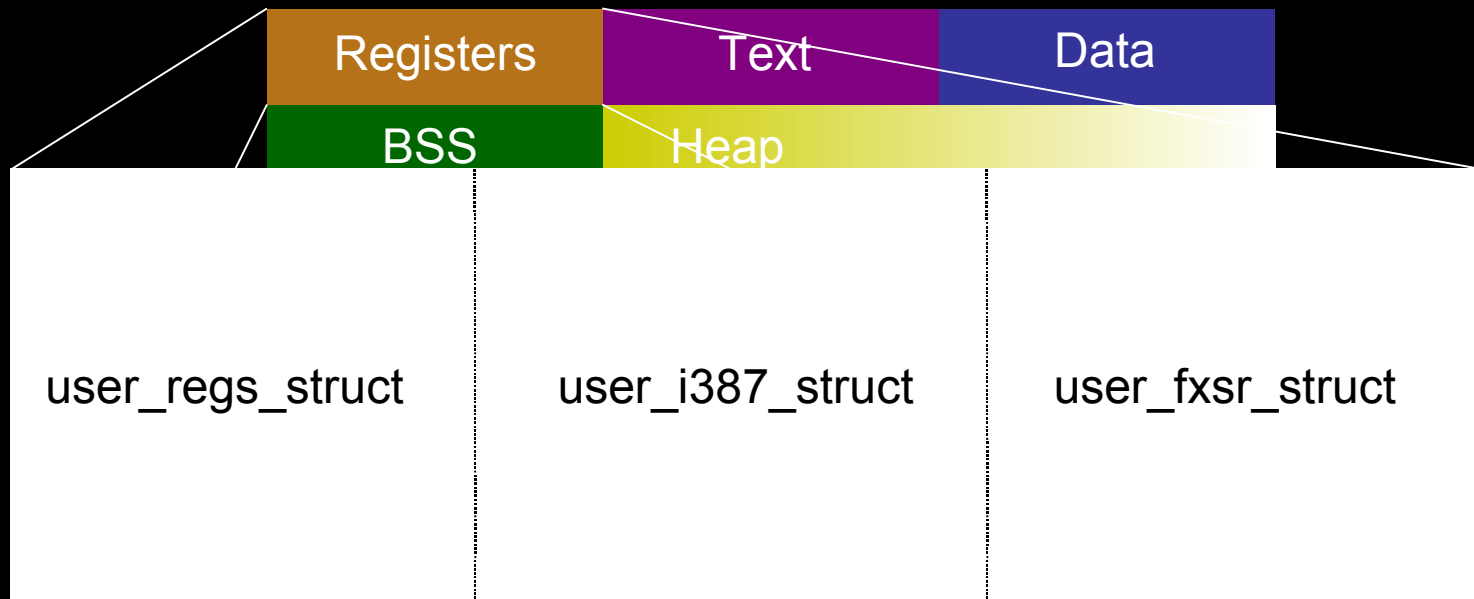
- Two files are produced by the freezing process \*.freeze & \*.frore.
- \*.freeze contains:
  - Register Data
  - Memory Segments
- \*.frore contains:
  - Memory Segment ranges & perms

```
1010001010100010110010010001010101
100010111000101001101001010101010
0001001001110001100001000100101
100010111001011*10101100
0110101110101001110101010101010
1001001110001001010
11000001010010100
10010010011100
```



```
1010001010100010110010010001010101
100010111000101001101001010101010
0001001001110001100001000100101
100010111001011*10101100
0110101110101001110101010101010
1001001110001001010
11000001010010100
10010010011100
```

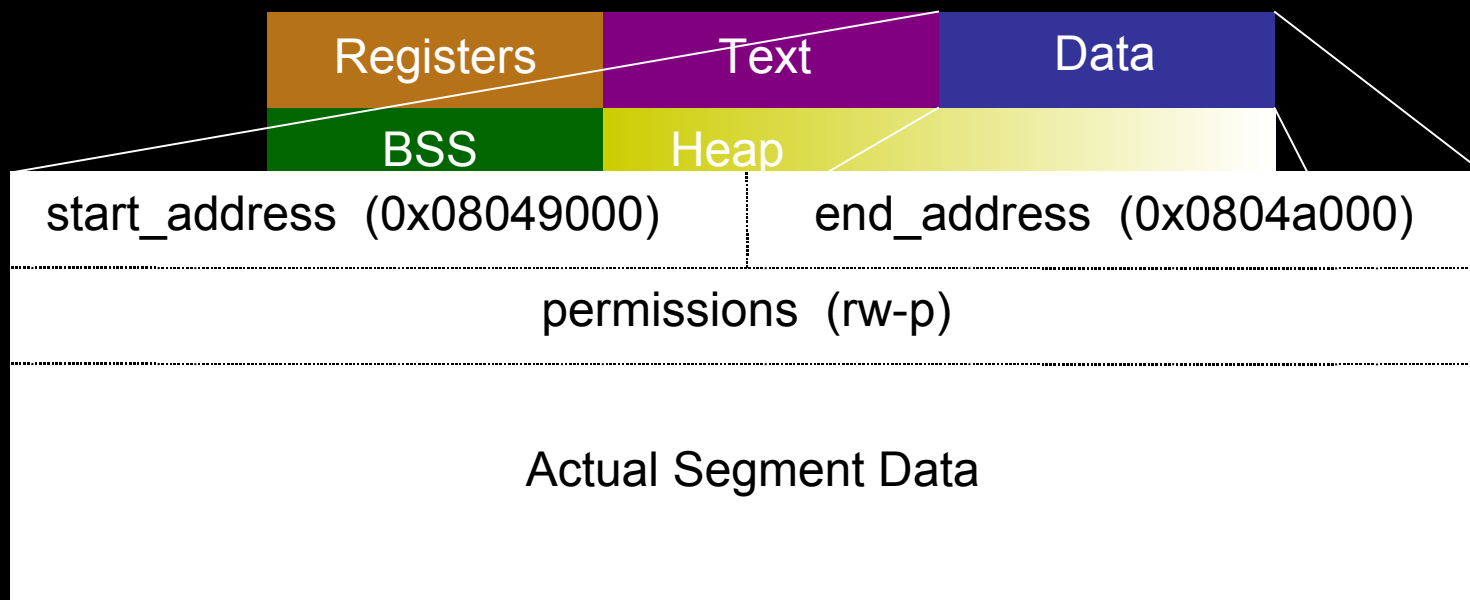
# \*.freeze In Depth(2)



ptrace() provides a way to retrieve these three structures, containing the process register data.

1010001010100010110010010001010101  
10001011000101001101001010101010  
0001001001110001100001000100101  
1000101110011101101110011100  
0110101110101011101001110011100  
1001001110001001010  
11000001010010100  
10010010011100

# \*.freeze In Depth(3)

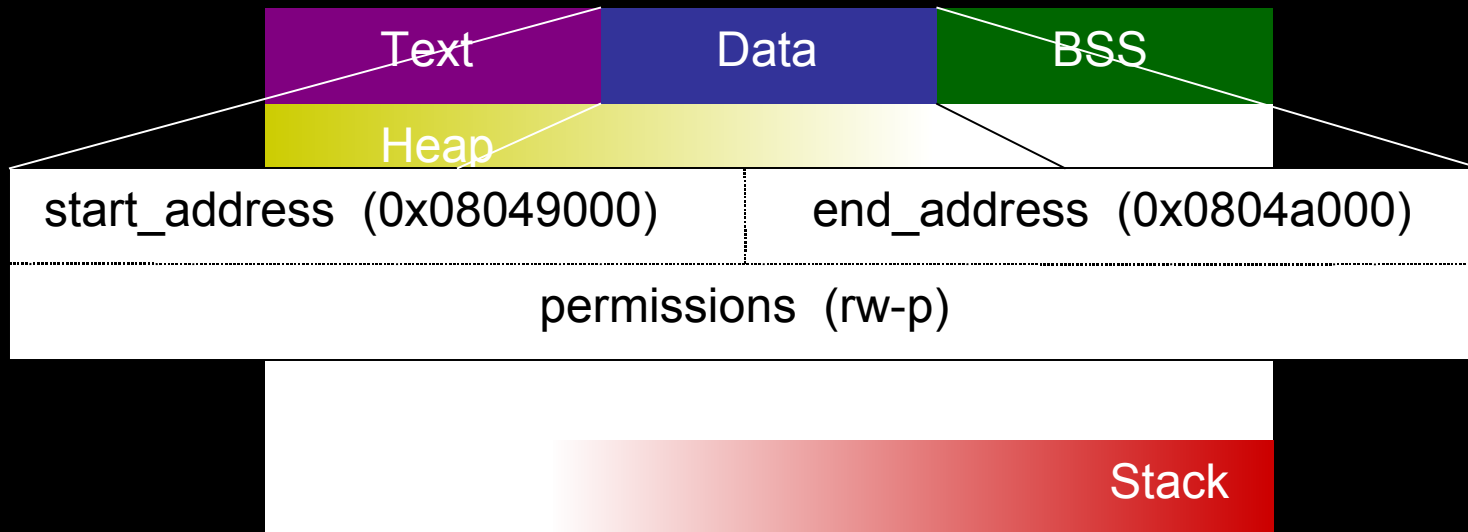


The remainder of the memory segments have much the same form, addresses followed by permissions and then the actual data of the segment.



1010001010100010110010010001010101  
10001011000101001101001010101010  
0001001001110001100001000100101  
1000101110010110101011100  
0110101110101001110100101010  
1001001110001001010  
11000001010010100  
10010010011100

# \*.frore In Depth



\*.frore is a template for the process. It contains everything \*.freeze contains except the registers and the actual segment data.

# Thoughts on Freezing

- Freezing a process and letting it slip away into digital oblivion is easy once we know how to get access to the necessary information.
- As intended, \*.freeze & \*.frore are portable across Linux systems with similar architecture (x86).
- Now, how do we bring it back to life...?

# The Big Thaw

- How do the files produced by freezing a process become a process of their own?
- In Linux, the only way to make a *new* process is with `fork()` or `clone()`.
- Replacing the context of a running program is the job of `exec()`, but it takes an executable image, not our cryo-files.

# A New Executable?

- Could a new executable be built, using the saved data, that would produce a running clone of our frozen program.
- Unfortunately no. An ELF or COFF executable cannot specify the size or contents of the stack, nor duplicate the new mappings of the heap at any given address.

# A New Executable? (2)

- Furthermore, if the program had a BSS segment, an executable cannot be set up to fill the contents of such a segment (the Kernel allocates clean pages for BSS).
- Also, as we see later signal handlers and file descriptors cannot be handled this way.

# A New Executable?(3)

- Perhaps a loader could be written with process revival specifically in mind. That is another project entirely, not implemented here.

# Seen “The Exorcist”?

- The idea is to “possess” an existing process, replacing its execution context with our process.
- What kind of process makes a good candidate for possession?
  - One with an *identical* memory mapping.
  - One with at *least* the same stack depth.
  - More features are needed, but these prove the point.



# The Mapping Conundrum

- Can a process change the memory mapping of another process? Not by any means... changes to a process's memory mapping must be voluntary or performed by the kernel.
- The ability to change a child's memory mapping would be a great addition to the `ptrace()` system call!!

# Voluntary Mapping

- Our “host” process has to change itself, via `mmap()` and other similar calls.
- The catch... How can a process shrink or expand its own text segment?
- It might possibly be done at user level, but not easily.
- *Divine* assistance is necessary... the Kernel.

# /dev/overseer

- A kernel module has the power to re-map a process on-the-fly.
- Using Kernel functions, a process could be remapped to serve as a suitable host for possession.
- The host process must simply write the contents of the \*.fcore file to the overseer module at it will be re-mapped.

# It's Getting Warmer...

- A new *library* function is created:

```
int thaw_process(int PID, int ops)
```

- Using the overseer module and `ptrace()`, a *new* process is created with the same execution context as the original frozen process.

- `freeze_process` & `thaw_process` are a big step but only a small part of migrating more complicated processes.

# What Escaped Freezing?

➤ Unfortunately this simple freeze does not handle saving such things as:

- open file descriptors
- installed signal handlers
- pending timers
- relationships to other processes
- many others...

# Enter OVERSEER

- Everything the simple freeze missed was installed via a system call (like `open()`, `signal()`, ...).
- In Linux, system calls can be intercepted with `ptrace()`.
- With `ptrace()`, the entry and exit of each system call can be inspected, then the appropriate action taken by the parent process.



# OVERSEER

- As its name suggests, the overseer will monitor every aspect of another process that is necessary for freezing.
- Overseer builds lists of installed signal handlers, open file descriptors (with their associated files) child processes, and some other items.



# Intercepting System Calls

- With `ptrace()`, delivery of a signal to the parent can be arranged when its child enters or exits a system call (the child gets stopped by the Kernel).
- The parent must look in the child's registers to see *which* system call caused the signal, and *what* its arguments were.
- Inspection of the registers is also necessary to determine the return value of the system call.

# System Calls in Linux

- System calls are requests by user processes for services from the Kernel.
- How is a system call actually made?
  - The appropriate syscall num is loaded into EAX.
  - The arguments are loaded into EBX, ECX, EDX, ...
  - Interrupt 0x80 is invoked, transferring control to the Kernel.

# \*.frost

- Information OVERSEER gathers throughout the monitoring process is written into a \*.frost file at the time of freezing.
- For \*.frost to be complete, the process must obviously be monitored from the beginning of its execution.

# OVERSEER (File Handling)

- Calls like `open()` and `dup()` cause new entries to the file descriptor (FD) list.
- Calls like `read()` and `lseek()` adjust the offset for the entries in the FD list.
- And so on...

# OVERSEER (Signal Handling)

- The installation of signal handlers is the passing of a certain struct to the Kernel. If intercepted and saved, the signal handler could be installed in the host later on.

# OVERSEER (Process Family)

- The `fork()` system call introduces an interesting twist to the overseer program.
- This migration implementation allows for communication between parent and child.
- When a `fork()` is encountered, the overseer forks as well, and its child monitors the forked process...




# OVERSEER (Process Family)

- When migration is not involved, parent/child communication is usually not a big deal.
- How can this communication continue on the destination machine?
  - PIDs are different.
  - calls such as getpid() and getppid() will NOT return the same values.
  - static storage of PIDs (by the processes) will be wrong.



```
1010001010100010110010010001010101
10001011000101001101001010101010
0001001001110001100001000100101
100010111001011010101100
0110101110101001110100101010
1001001110001001010
11000001010010100
10010010011100
```

# \*.frore In Depth



\*.frore contains the structures and other things necessary to get these “services” reinitialized.

# Demo Time

- An example process “theprog” is going to be migrated.
- Add a fork() or two...
- Look at some code maybe...
- Some requests maybe...

```
10010001010100010110010010001010101
011010010101010101000
001000100101110
110110101110010101010
101001110100101010100010
001001010101011100101001
0000101001010001001011010101010
01001001110001010010001011100010
```

# Future Expansion

- There is a lot that even the mighty OVERSEER does not handle quite yet.
  - Multi-threaded processes.
  - Shared memory (in the same family).
  - Pipes/FIFOs (in the same family).
  - Sockets (set up a proxy... ugly).

# Any Questions?

Please be gentle!!!

# The End

Now migrate yourself home...