

Processes

Overview

- ▶ Threads vs Processes
- ▶ Clients
- ▶ Servers
- ▶ Virtualization
- ▶ Code Migration

Threads versus Processes

Threads:

- ▶ Improve application responsiveness by allowing them to not block

Threads versus Processes

Threads:

- ▶ Improve application responsiveness by allowing them to not block
- ▶ Allows for parallel computation resulting in higher speed on many-core systems

Threads versus Processes

Threads:

- ▶ Improve application responsiveness by allowing them to not block
- ▶ Allows for parallel computation resulting in higher speed on many-core systems
- ▶ Easier to structure many applications as a collection of cooperating threads

Threads versus Processes

Threads:

- ▶ Improve application responsiveness by allowing them to not block
- ▶ Allows for parallel computation resulting in higher speed on many-core systems
- ▶ Easier to structure many applications as a collection of cooperating threads
- ▶ Higher performance compared to multiple processes since switching between threads takes less time

Threads versus Processes

Threads:

- ▶ Improve application responsiveness by allowing them to not block
- ▶ Allows for parallel computation resulting in higher speed on many-core systems
- ▶ Easier to structure many applications as a collection of cooperating threads
- ▶ Higher performance compared to multiple processes since switching between threads takes less time

Processes:

- ▶ Simpler to synchronize as processes have separate memory space

Threads versus Processes

Threads:

- ▶ Improve application responsiveness by allowing them to not block
- ▶ Allows for parallel computation resulting in higher speed on many-core systems
- ▶ Easier to structure many applications as a collection of cooperating threads
- ▶ Higher performance compared to multiple processes since switching between threads takes less time

Processes:

- ▶ Simpler to synchronize as processes have separate memory space
- ▶ Allows for parallel computation resulting in higher speed on many-core systems

Thread Implementations

- ▶ User-space threads:

Thread Implementations

- ▶ **User-space threads:**
 - ▶ Creating/destroying threads is inexpensive

Thread Implementations

- ▶ **User-space threads:**
 - ▶ Creating/destroying threads is inexpensive
 - ▶ Switching context is very fast

Thread Implementations

- ▶ **User-space threads:**
 - ▶ Creating/destroying threads is inexpensive
 - ▶ Switching context is very fast
 - ▶ But invocation of a blocking system call blocks all threads...

Thread Implementations

- ▶ **User-space threads:**
 - ▶ Creating/destroying threads is inexpensive
 - ▶ Switching context is very fast
 - ▶ But invocation of a blocking system call blocks all threads...
 - ▶ Cannot make use of multiple cores

Thread Implementations

- ▶ **User-space threads:**
 - ▶ Creating/destroying threads is inexpensive
 - ▶ Switching context is very fast
 - ▶ But invocation of a blocking system call blocks all threads...
 - ▶ Cannot make use of multiple cores
- ▶ **Kernel threads:**

Thread Implementations

- ▶ **User-space threads:**
 - ▶ Creating/destroying threads is inexpensive
 - ▶ Switching context is very fast
 - ▶ But invocation of a blocking system call blocks all threads...
 - ▶ Cannot make use of multiple cores
- ▶ **Kernel threads:**
 - ▶ Overcomes the last two issues with user-space threads but loses performance

Thread Implementations

- ▶ **User-space threads:**
 - ▶ Creating/destroying threads is inexpensive
 - ▶ Switching context is very fast
 - ▶ But invocation of a blocking system call blocks all threads...
 - ▶ Cannot make use of multiple cores
- ▶ **Kernel threads:**
 - ▶ Overcomes the last two issues with user-space threads but loses performance
- ▶ **Light-Weight Processes (LWP):** Hybrid model. Multiple LWP/threads run inside a single (heavy-weight) process. In addition, the system offers a user-level threads package.

- ▶ A **server** is a process implementing a specific service on behalf of a collection of clients

- ▶ A **server** is a process implementing a specific service on behalf of a collection of clients
- ▶ A server can be **iterative** or **concurrent**. Concurrent servers can be *multi-threaded* or *multi-process*.

- ▶ A **server** is a process implementing a specific service on behalf of a collection of clients
- ▶ A server can be **iterative** or **concurrent**. Concurrent servers can be *multi-threaded* or *multi-process*.
- ▶ Characteristics of server implementations.

Model	Characteristics
Single-threaded	No parallelism, blocking system calls
Multi-threaded	Parallelism, blocking system calls
Finite state machine	Parallelism, nonblocking system calls

Design Issues for Servers (2)

- ▶ How does a client find a server? Need to know the end-point or port and the host address.

Design Issues for Servers (2)

- ▶ How does a client find a server? Need to know the end-point or port and the host address.
 - ▶ Statically assigned like well known servers like HTTP on port 80.

Design Issues for Servers (2)

- ▶ How does a client find a server? Need to know the end-point or port and the host address.
 - ▶ Statically assigned like well known servers like HTTP on port 80.
 - ▶ Look-up service provided by a special directory server.

Design Issues for Servers (2)

- ▶ How does a client find a server? Need to know the end-point or port and the host address.
 - ▶ Statically assigned like well known servers like HTTP on port 80.
 - ▶ Look-up service provided by a special directory server.
 - ▶ Using a superserver that selects on multiple ports and forks off the appropriate server when a request comes in.